

Contour — the measurement picks the structure

`vrp_ratio < 1.30` → **NO_TRADE** implied is not rich enough to sell

`skew_z ≥ +0.8` → **PUT_CREDIT_SPREAD** puts rich – sell puts, not cheap calls

`skew_z ≤ -0.8` → **CALL_CREDIT_SPREAD** calls rich – sell calls, not cheap puts

`otherwise` → **IRON_CONDOR** both sides fair – sell both

AI LOGIC — THE LEASH

Every wired model output can only make the agent trade **less**. Not a promise — a property of the import graph.

THE MODEL MAY

- Name event windows to stand down in
- Veto a proposed structure
- Stand the whole book down

THE MODEL MAY NEVER

- Choose a strike
- **Size** or price a position
- Reverse or widen anything

ALPACA INFRASTRUCTURE

- **Orders go through the CLI, and that is a finding.** The MCP server cannot place multi-leg orders — the `legs` array arrives as a JSON string and fails validation. `alpaca-mcp-server` #97, open since July. The CLI places the identical 4-leg order correctly; our fix is upstream as #118.
- Reads are `alpaca-py`: chain snapshots carry Greeks but no `open_interest`, so they are merged with Trading API contract objects.
- Entries are a **three-rung limit ladder** from mid toward the bid, never a market order. `reconcile()` reads actual per-leg fills: paper issues random partials, and trusting the request puts legs out of ratio.
- The journal is an **append-only SHA-256 hash chain**. The dashboard recomputes it in your browser and prints the same verdict the CLI does.
- Autonomy is GitHub Actions: pre-open planning, a cycle every 15 minutes, a scheduled Thursday flatten. **No `pull_request` trigger** on the credentialed workflow.
- **No resting stop exists on a multi-leg position** — so the options book polls. A single equity leg *can* rest one, and the directional sleeve parks its stop GTC at the broker: the only exit that works overnight.

MARKET ANALYSIS

The volatility premium is real, and it is **not uniform**. That dispersion is the entire opportunity.

SPY

1.42

paid to sell

QQQ

1.17

not paid enough

IWM

1.20

not paid enough

- Short premium is the crowded retail options trade — and the crowd sells **one structure regardless of the surface**.
- **Friction decides the universe, not opinion.** Single-name weeklies cost \$40–80 round trip against a \$30–42 modelled edge — it loses to costs before it starts. Three ETFs cost \$8–20.
- Alpaca serves **no earnings-date endpoint on any plan**, so a single-name design hangs its most important gate on data that does not exist.

COMPETITIVE ANALYSIS

SUBMISSION	ITS OWN DESCRIPTION
Horizon Blackline	LLM proposes, deterministic risk gates authorize, hash-chained and auditable
VRP Engine	Harvests the variance risk premium with defined-risk spreads and risk gates
AEGIS-Q	Bounded AI selects a pre-validated bullish or bearish spread — or abstains
EdgeStack	Journals every trade and every refusal
Contour	Chooses which of four structures to sell, from measured 25-delta skew. And is runnable without our credentials.

“LLM proposes, gates dispose, everything journaled” is table stakes now

```
python -m contour --replay
```

PERFORMANCE, ATTRIBUTED

The criterion asks for the P&L of **the submitted agent**. This account holds two traders, and the broker records which is which. 2026-09-04 07:40 UTC

PLACED BY	P&L	OF START NAV
The agent — every id it chose	+\$147.55	+0.15%
The operator — three discretionary tail trades	-\$642.00	-0.64%
Account total	-\$494.45	-0.49%

Not our bookkeeping — a field the broker stamps. Entries are prefixed contour- by loop.py; exits are named after the entry they close. Nothing else in the account is. A symbol touched by both is charged to the operator — the published number is the pessimistic one.

REVENUE MODEL

The audit layer is the product

The gate engine and the hash-chained decision record, licensed to brokers and RIAs who have to **defend** an automated decision after the fact. That is the durable asset here — not the alpha, which decays.

Own capital

P&L is the revenue and no registration question arises. The honest ceiling on defined-risk premium selling is the credit, so this scales with capital, not with claims — and where we stepped outside that ceiling, the write-up names the trade and its negative expectancy.

Not signal subscriptions. Selling trade recommendations is investment advice and needs registration — saying so out loud is a credibility position, not a limitation.

ROADMAP

SHIPPED

- 3 ETFs, one locked expiry, 15-minute cycle
- **A \$30k long-XXX sleeve** beside the options book — variance, not edge — funded *out of* the same -4% floor, not beside it
- **A \$4.4k long-call tail**, added *and closed* 2026-09-01 — the one negative-EV trade here, labelled as such. Sold for \$3,113, realising **-\$1,320**: the entire drawdown. A 387-cycle backtest found no edge worth paying ~15% over fair value to lever
- **A \$1.1k TXXX call tail**, placed on instruction after the evidence against it was recorded — a gap-down bounce tests as noise ($t = +0.42$). It breaks no gate, but used the last room in front of the floor: the entry ramp is *closed to zero* for the rest of the contest
- Twelve gates plus seven, 298 tests, chain verified in CI
- Record / replay: a fixture reruns the pipeline with no credentials
- Sizing from three published trend systems, after an audit caught the model anchoring at 0.5 for sixteen straight cycles

NEXT

- **Trend-aware structure selection.** The map breaks ties toward the condor — which sells calls into a confirmed uptrend. Three trend systems say that is the wrong default; skew alone should not decide
- Expiry laddering and rolls
- Skew priors **learned** per underlying instead of hard-coded
- **More backtest history.** The harness is built and has run — 387 cycles against real historical option prices, importing the agent's own selection and gate code. Alpaca's history starts 2024-01-18, so 2.5 years is the entire available sample
- Portfolio vega and gamma caps, not per-position max loss alone
- Paid feed to close the indicative-vs-NBBO gap